

Figure 5 — Memory Pareto frontier: vorn methods cluster ~20-22 GB peak allocated;
attention-weight methods cluster ~29-58 GB. Same A100-80GB profile, dramatically different memory cost.

